

AI BUILDERS GUILD

7:30 Spark Session Member Notes

10 Deep-Dive Threads for Builders

September 4, 2026 · Member notes with source links

AT A GLANCE

- Deeper cuts on the coffee threads
- Starter + keep-going + clickable sources
- Drop source URLs in chat as we go
- Leftover heat returns to morning coffee hour
- Bootcamp stays L1 GitHub → L2 Netlify → L3 VS Code
- Oath: Never quit. Every hand you lift lifts your own.

1. Fable 5.1 vs Mythos 5.1 — Same Model, Different Gates

STARTER

Anthropic is explicit: Fable 5.1 is the generally available Mythos-class model; Mythos 5.1 is the same capability with cyber and biology safeguards lifted, offered only to approved Project Glasswing customers. Pricing listed for Mythos starts at \$10/\$50, matching Fable. Both are Covered Models with 30-day retention. Platform docs warn agent loops may see one tool call per turn where Fable 5 batched several. Spark is a policy lab. If you cannot get Mythos, what do you actually change in Fable — model, tools, or the jobs you refuse to run.

KEEP GOING

- Screen-share your model picker — Fable 5.1, Opus 5, or something else?
- What task would you not run on Fable because the fallback might fire?
- Would you accept 30-day retention on a Guild repo, or switch providers?

SOURCES (clickable)

- <https://www.anthropic.com/claude-fable-and-mythos-5-1>
- <https://www.anthropic.com/claude/mythos>
- <https://platform.claude.com/docs/en/models/fable-5-1/whats-new-fable-5-1>
- <https://www.anthropic.com/claude-fable-5-1-mythos-5-1-system-card>

2. Route GPT-5.6 Like a Team, Not a Chatbot

STARTER

OpenAI's GPT-5.6 page splits work across Sol, Terra, and Luna, then adds effort controls. Max gives more time than xhigh. Ultra is called out for Pro and Enterprise in ChatGPT Work. Sol is \$5/\$30. Terra is \$2.50/\$15. Luna is \$1/\$6. The quote pile is coding agents, decks, and frontend. Spark wants a two-axis card you can steal: model tier on one axis, effort on the other. Draw it for one real Guild job — a PR, a landing page, or a research brief — and say where you would spend.

KEEP GOING

- Pick one live task. Sol, Terra, or Luna — and which effort?
 - Where has high effort wasted money for you this month?
 - Would you let an agent pick the tier, or lock it in the harness?
-

SOURCES (clickable)

- <https://openai.com/index/gpt-5-6/>
- <https://openai.com/index/gpt-5-6/>
- <https://cursor.com/changelog>

3. Grok 4.6 Evals vs Your Repo

STARTER

xAI published Grok 4.6 High next to Grok 4.5, GPT-5.6 Sol Max, and Fable 5 Max. On their table, 4.6 ties Sol at 61 on AA Intelligence, sits near Fable on GDPVal-AA, and trails Sol and Fable on Terminal-Bench. Price is the wedge: \$2/\$6 versus Fable's \$10/\$50 and Sol's \$5/\$30. Spark is not a screenshot of a chart. It is a 15-minute bake-off on one of your repos. Same prompt, two models, ship the diff that actually compiles.

KEEP GOING

- Which one eval on that chart matches work you actually do?
 - Run the same prompt on 4.6 and Fable or Sol — who wins on your files?
 - If 4.6 is close enough, what do you do with the money you save?
-

SOURCES (clickable)

- <https://x.ai/news/grok-4-6>
- <https://x.ai/build>
- <https://www.anthropic.com/claude-fable-and-mythos-5-1>
- <https://openai.com/index/gpt-5-6/>

4. Self-Hosted Cursor — What Leaves the Network

STARTER

Cursor's self-hosted docs split the loop. Cursor still runs inference, planning, and the agent loop. Your worker runs the tools. Personal machines use a user login or API key. Team pools need Enterprise plus a service account key. Admins can allow or require self-hosted. You can connect up to 200 workers per user. Cursor-managed Cloud Agents remain the default for most teams. Spark is a whiteboard: draw what stays on the worker, what goes to Cursor, and what a client would not sign.

KEEP GOING

- Draw the trust boundary — model in Cursor's cloud, shell on your box. Acceptable?
 - Personal My Machines, or only a team pool on a VM?
 - Require self-hosted for client work, or keep Cursor-managed for speed?
-

SOURCES (clickable)

- <https://cursor.com/docs/cloud-agent/self-hosted>
- <https://cursor.com/changelog>
- <https://cursor.com/docs/cloud-agent/self-hosted/my-machines>
- <https://cursor.com/docs/cloud-agent/self-hosted/choose-runtime>

5. Computer Use on a Worker You Own

STARTER

The same September 2 Cursor drop adds computer use on Linux and Mac self-hosted workers. On a Mac, the CLI installs a Cursor Computer Use helper and needs Accessibility plus Screen Recording. On Linux, the worker wants a desktop stack so the agent can click, type, screenshot, and drive a browser. You can watch the desktop or take control from Cursor. Spark is a demo-or-refuse. Who would grant those permissions on a Guild Mac mini, and who would only allow it inside a disposable VM.

KEEP GOING

- Mac mini on the desk, or a throwaway Linux VM — which gets computer use?
 - What app would you refuse to let an agent click through?
 - Watch the desktop live, or only read the log after?
-

SOURCES (clickable)

- <https://cursor.com/changelog>
- <https://cursor.com/docs/cloud-agent/self-hosted>
- <https://cursor.com/docs/cloud-agent/self-hosted/computer-use>
- <https://cursor.com/docs/cloud-agent/self-hosted/integrations>

6. Should Copilot Be Allowed to Approve

STARTER

GitHub now lets Copilot code review approve a pull request. Skills and MCP already feed that review. MCP calls stay read-only. Comments can show when a skill or MCP server wrote them. That is a new merge policy, not a party trick. Spark wants a written rule you could paste into a Guild repo tonight: which paths Copilot may approve, which paths always need a human, and whether Fable 5.1 in Copilot changes that rule.

KEEP GOING

- Write the rule in one sentence — Copilot may approve X, never Y.
 - Would you turn on Copilot approve for Bootcamp student repos?
 - Does Fable 5.1 in Copilot make approval safer, or just more confident?
-

SOURCES (clickable)

- <https://github.blog/changelog/2026-09-01-copilot-code-review-can-now-approve-pull-requests>
- <https://github.blog/changelog/2026-07-29-copilot-code-review-agent-skills-and-mcp-now-generally-available>
- <https://docs.github.com/copilot/concepts/agents/about-agent-skills>
- <https://github.blog/changelog/2026-09-01-claude-fable-5-1-generally-available-in-github-copilot>

7. Skills vs MCP — Pick a Lane for Review

STARTER

GitHub's agent-skills docs and the July 29 changelog draw a clean split. A SKILL.md under .github/skills teaches Copilot how your team reviews. An MCP server pulls live context from issue trackers and docs. Both now show up as attribution on review comments. MCP in review is read-only. Cursor and VS Code have their own MCP clients, so the same server can appear in three places with three permission models. Spark is a packing list. One skill, one MCP, one repo — ship it this week or say why not.

KEEP GOING

- Would you write a SKILL.md for Guild PRs before you add another MCP server?
 - Which third-party MCP is worth read-only review context?
 - Same MCP in Cursor, VS Code, and Copilot review — or keep them separate?
-

SOURCES (clickable)

- <https://docs.github.com/copilot/concepts/agents/about-agent-skills>
- <https://github.blog/changelog/2026-07-29-copilot-code-review-agent-skills-and-mcp-now-generally-available>
- <https://modelcontextprotocol.io/>
- <https://cursor.com/docs/context/mcp>

8. Model Hardware Standard — Agents Touch the Physical World

STARTER

August 27: Anthropic opened a research preview of the Model Hardware Standard, a shared spec for AI agents to operate physical devices with a first group of labs and manufacturers. The Fable 5.1 post points at it. This is not a Guild lab protocol and not a how-to. It is a product-shape question. Software agents are growing a hardware USB-C. Spark stays high-level: where would you even allow an agent near a device you own, and where is that a hard no.

KEEP GOING

- Printer, 3D printer, or never — where does an agent get physical reach?
 - Would you want a standard like MCP for hardware, or one vendor's stack?
 - What has to be true before you would let an agent touch a machine at work?
-

SOURCES (clickable)

- <https://www.anthropic.com/news/model-hardware-standard-research-preview>
- <https://www.anthropic.com/claude-fable-and-mythos-5-1>
- <https://www.anthropic.com/news>

9. Origin, GitHub, or a Worker — Where Does the Code Live

STARTER

Cursor Origin can host the repo. GitHub can stay source of truth with Origin as a mirror. Self-hosted workers can run the tools without giving Cursor the shell. VS Code can open Agents without a GitHub login if you bring a Claude API key. That is four different homes for the same project. Spark is a map for one Guild app: Bootcamp still teaches GitHub as L1. Say out loud whether this week's stack still starts there, or whether Origin and self-hosted workers changed the first click.

KEEP GOING

- For a new Guild app tonight: GitHub first, Origin first, or local only?
 - What never goes on Origin — client code, secrets, or everything until a human clones?
 - Does Bootcamp L1 still start on github.com, or do we show Origin too?
-

SOURCES (clickable)

- <https://cursor.com/docs/origin>
- <https://cursor.com/changelog>
- <https://cursor.com/docs/cloud-agent/self-hosted>
- <https://github.blog/changelog/2026-09-04-github-copilot-in-vs-code-august-2026-releases>

10. Guild Rhythm — Calendar, Bootcamp, Oath

STARTER

Spark is the overflow lane, not a second keynote. Confirm times on the SKOOL Calendar tab. Unfinished coffee threads live here; leftover Spark heat goes back to morning coffee hour. Bootcamp remains L1 GitHub, L2 Netflix, L3 VS Code until the pack is extended. Close on the oath and send people out building.

KEEP GOING

- Which coffee thread deserves the next twenty minutes on a screen-share?
 - Newcomers: Calendar tab — which meeting is realistic this week?
 - One hand you will lift before the next coffee hour.
-

SOURCES (clickable)

- <https://www.skool.com/ai-builders-guild-9932>

Guild Reminder

Six weekly meetings live in the **Calendar** tab on SKOOL. Coffee Hour seeds threads; **7:30 Spark** goes deeper. Bootcamp path: **L1 GitHub** → **L2 Netlify** → **L3 VS Code**.

"Never quit. Every hand you lift lifts your own."

— AI Builders Guild